



NORTHEASTERN UNIVERSITY, KHOURY COLLEGE OF COMPUTER SCIENCE

CS 6220 Data Mining — Assignment 5

Due: February 1, 2023(100 points)

YOUR NAME
YOUR GIT USERNAME
YOUR E-MAIL

Naïve Bayes, Bayes Rules

The original performance of [acoustic classification for Parkinsons Disease](#) leverages speech recordings from controlled subject responses from variety of questions. The task in the competition was to detect whether or not a person X had Parkinsons disease from a sampling of data. As of 2018, the state of the art classifiers have achieved 90% correct classification on a held out dataset, both for subjects who had Parkinsons and those who did not (at equal rates). So, when classifier Y sees person X , it works correctly 90% of the time.

Question 1 [25 pts total]

Let's say that we run a clinic. This clinic leverages this classifier, which has 90% accuracy. Also, let us say that we know that our current patient load is that 10% of the population have Parkinsons and 90% of the population do not. Let's also say that we're seeing patient X , and the classification algorithm has detected that they have Parkinson's disease. What's the probability that indeed X has Parkinson's disease?

Come up with the numerical solution, and show your written work.

Gradient Descent - Logistic Regression

We'll explore logistic regression with some data from Kaggle today. If you've done the lab in class, then this simply becomes an exercise of data transformations and preprocessing. (Since there is only one label that you are predicting, you may also have to ensure that the dimensions

are correct.)

Recall that the logistic regression with binary cross-entropy cost function appears in the following forms:

$$\mathcal{L}(W, b) = - \sum_i y_i \log h_{W,b}(\mathbf{x}) + (1 - y_i) \log (1 - h_{W,b}(\mathbf{x}))$$

where

$$h_{W,b}(\mathbf{x}) = \sigma(W^T \mathbf{x} + b)$$

and

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$

Also, recall that you can solve the logistic regression problem entirely through the gradients as provided below:

- The gradient of \mathcal{L} with respect to W : $\nabla_W \mathcal{L} = (\mathbf{y} - h_{W,b}(\mathbf{x})) \cdot \mathbf{x}^T$
- The gradient of \mathcal{L} with respect to b : $\nabla_b \mathcal{L} = (\mathbf{y} - h_{W,b}(\mathbf{x}))$

Question 2a.)

Prove that the derivative of the sigmoid function is $\nabla_z \sigma(z) = \sigma(z) (1 - \sigma(z))$.

Question 2b.)

The sinking of the Titanic is one of the most infamous shipwrecks in history. On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there were not enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew. While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others. In this challenge, Kaggle asks you to build a predictive model that answers the question: “what sorts of people were more likely to survive?” using passenger data (i.e., name, age, gender, socio-economic class, etc).

Some tips and tricks:

- For features with multiple values, try making one-hot encoded data. For example, number of siblings, you can try to have a four column array.

No Siblings	1 Sibling	2 Siblings	3 or more Siblings
0	1	0	0

- Play around with your learning rate. If your learning is too high, then you won’t converge to the right answer. I used X features with a learning rate of.

- You needn't use all the features. Just because they're there, doesn't mean they're always useful.

See if you can do better than I can with a logistic regression. On the test set, I reached 79.2% accuracy at a threshold of 0.5.

Question 3)